



# Long Short-Term Memory Networks



# 1. Why LSTMs?

What is wrong with me? :( asks RNN




# Why LSTMs?

- Traditional RNNs had a lot of matrix multiplication involved.
- Naturally this caused the vanishing gradients problem during back-propagation.



# Vanishing Gradients in RNNs



$$\frac{\partial J^{(4)}}{\partial h^{(1)}} = \frac{\partial h^{(2)}}{\partial h^{(1)}} \times$$

$$\frac{\partial h^{(3)}}{\partial h^{(2)}} \times$$

$$\frac{\partial h^{(4)}}{\partial h^{(3)}} \times \frac{\partial J^{(4)}}{\partial h^{(4)}}$$

# Why is vanishing gradients a problem?

- Gradient signal from far away is lost because it's much smaller than the gradient from close-by.
- As a result, model weights are updated only with respect to near effects, not long-term effects.

# Why is vanishing gradients a problem?

- Suppose a language modelling task:

The leaves of the Banyan tree \_\_\_\_ ... (is or are?).

Correct answer: The leaves of the Banyan tree are ...

# Why is vanishing gradients a problem?

- An RNN suffering from vanishing gradients may answer by giving weightage to sequential recency.
- **Syntactic Recency:** The **leaves** of the Banyan tree are ...
- **Sequential Recency:** The leaves of the Banyan tree is ...

# How to fix the vanishing gradients problem?

- In an RNN, the hidden state  $h^{(t)}$  is constantly being rewritten. Hence, it is very difficult for the model to preserve information over many time steps.
- LSTMs propose a separate memory for the RNNs.



# How does an LSTM solve vanishing gradients?


- LSTMs reduce the no. of matrix multiplications in favour of Hadamard product and addition of Matrices.
- LSTMs make it easier for the RNN to preserve information over many time steps.



## 2. What are Vanilla LSTMs?

Interesting stuff ahead...

# Vanilla LSTMs



# Vanilla LSTMs Forward Pass

$$\bar{\mathbf{z}}^t = \mathbf{W}_z \mathbf{x}^t + \mathbf{R}_z \mathbf{y}^{t-1} + \mathbf{b}_z$$

$$\mathbf{z}^t = g(\bar{\mathbf{z}}^t)$$

*block input*

$$\bar{\mathbf{i}}^t = \mathbf{W}_i \mathbf{x}^t + \mathbf{R}_i \mathbf{y}^{t-1} + \mathbf{p}_i \odot \mathbf{c}^{t-1} + \mathbf{b}_i$$

$$\mathbf{i}^t = \sigma(\bar{\mathbf{i}}^t)$$

*input gate*

$$\bar{\mathbf{f}}^t = \mathbf{W}_f \mathbf{x}^t + \mathbf{R}_f \mathbf{y}^{t-1} + \mathbf{p}_f \odot \mathbf{c}^{t-1} + \mathbf{b}_f$$

$$\mathbf{f}^t = \sigma(\bar{\mathbf{f}}^t)$$

*forget gate*

$$\mathbf{c}^t = \mathbf{z}^t \odot \mathbf{i}^t + \mathbf{c}^{t-1} \odot \mathbf{f}^t$$

*cell*

$$\bar{\mathbf{o}}^t = \mathbf{W}_o \mathbf{x}^t + \mathbf{R}_o \mathbf{y}^{t-1} + \mathbf{p}_o \odot \mathbf{c}^t + \mathbf{b}_o$$

$$\mathbf{o}^t = \sigma(\bar{\mathbf{o}}^t)$$

*output gate*

$$\mathbf{y}^t = h(\mathbf{c}^t) \odot \mathbf{o}^t$$

*block output*

# Vanilla LSTMs Back-propagation through Time

$$\delta \mathbf{y}^t = \Delta^t + \mathbf{R}_z^T \delta \mathbf{z}^{t+1} + \mathbf{R}_i^T \delta \mathbf{i}^{t+1} + \mathbf{R}_f^T \delta \mathbf{f}^{t+1} + \mathbf{R}_o^T \delta \mathbf{o}^{t+1}$$

$$\delta \bar{\mathbf{o}}^t = \delta \mathbf{y}^t \odot h(\mathbf{c}^t) \odot \sigma'(\bar{\mathbf{o}}^t)$$

$$\begin{aligned} \delta \mathbf{c}^t &= \delta \mathbf{y}^t \odot \mathbf{o}^t \odot h'(\mathbf{c}^t) + \mathbf{p}_o \odot \delta \bar{\mathbf{o}}^t + \mathbf{p}_i \odot \delta \bar{\mathbf{i}}^{t+1} \\ &\quad + \mathbf{p}_f \odot \delta \bar{\mathbf{f}}^{t+1} + \delta \mathbf{c}^{t+1} \odot \mathbf{f}^{t+1} \end{aligned}$$

$$\delta \bar{\mathbf{f}}^t = \delta \mathbf{c}^t \odot \mathbf{c}^{t-1} \odot \sigma'(\bar{\mathbf{f}}^t)$$

$$\delta \bar{\mathbf{i}}^t = \delta \mathbf{c}^t \odot \mathbf{z}^t \odot \sigma'(\bar{\mathbf{i}}^t)$$

$$\delta \bar{\mathbf{z}}^t = \delta \mathbf{c}^t \odot \mathbf{i}^t \odot g'(\bar{\mathbf{z}}^t)$$



# 3. History of LSTMs

Crash Course: Modern History...



# Original LSTM

- The original LSTM did not have the forget gate and the peephole connections.
- The model was trained using a mixture of Real Time Recurrent Learning and Back-propagation through Time.
- It used *full gate recurrence*.

# Forget Gate

- Forget Gate was introduced in 1999, 4 years after the introduction of LSTMs in 1995.
- This allowed learning of continual tasks such as embedded Reber Grammar.



# Peephole Connections

- Peephole Connections were introduced in 2000, 5 years after the introduction of LSTMs in 1995.
- This allowed the cell to control the gates in order to learn precise timings.
- Output activation was omitted.

# Full Gradient

- Full Back-propagation through Time was used to train the LSTMs.
- LSTM gradients could be checked using finite differences, making practical implementations more reliable.



## 4. Evaluation Setup

On your mark, get set!...

# Setup

- We have to compare various variants of LSTM. So, we compare them across three different tasks on three different datasets.
- Hyper-parameters are separately tuned for each variant, so as to compare their best performances.

# Dataset Setup

- Each Dataset is split into Training, Validation and Testing Dataset.
- **TIMIT Speech Corpus:** *Speech Recognition*
- **IAM Online:** *Handwriting Recognition*
- **JSB Chorales:** *Music Modelling*

# Network Architecture

- Single LSTM hidden layer and sigmoid output layer for **JSB Chorales** task.
- Bidirectional LSTM for the other two tasks.
- Cross-Entropy Loss for **TIMIT** and **JSB Chorales** tasks.
- Connectionist Temporal Classification (CTC) Loss for **IAM Online**.

# Training

- Initial weights for all the networks were drawn from a normal distribution of s.d. = 0.1
- **Optimiser:** SGD + Nesterov Momentum
- The Learning Rate was rescaled by a factor of (1-momentum)
- Gradients were computed using BPTT
- 15 epochs

# LSTM Variants

- NIG: No Input Gate:  $i^{(t)} = 1$
- NFG: No Forget Gate:  $f^{(t)} = 1$
- NOG: No Output Gate:  $o^{(t)} = 1$
- NIAF: No Input Activation Function:  $g(x) = x$
- NOAF: No Output Activation Function:  $h(x) = x$
- CIFG: Coupled Input and Forget Gate:  $f^{(t)} = 1 - i^{(t)}$
- FGR: Full Gate Recurrence
- NP: No peepholes




## Full Gate Recurrence

$$\begin{aligned}\bar{\mathbf{i}}^t &= \mathbf{W}_i \mathbf{x}^t + \mathbf{R}_i \mathbf{y}^{t-1} + \mathbf{p}_i \odot \mathbf{c}^{t-1} + \mathbf{b}_i \\ &\quad + \mathbf{R}_{ii} \mathbf{i}^{t-1} + \mathbf{R}_{fi} \mathbf{f}^{t-1} + \mathbf{R}_{oi} \mathbf{o}^{t-1} \\ \bar{\mathbf{f}}^t &= \mathbf{W}_f \mathbf{x}^t + \mathbf{R}_f \mathbf{y}^{t-1} + \mathbf{p}_f \odot \mathbf{c}^{t-1} + \mathbf{b}_f \\ &\quad + \mathbf{R}_{if} \mathbf{i}^{t-1} + \mathbf{R}_{ff} \mathbf{f}^{t-1} + \mathbf{R}_{of} \mathbf{o}^{t-1} \\ \bar{\mathbf{o}}^t &= \mathbf{W}_o \mathbf{x}^t + \mathbf{R}_o \mathbf{y}^{t-1} + \mathbf{p}_o \odot \mathbf{c}^{t-1} + \mathbf{b}_o \\ &\quad + \mathbf{R}_{io} \mathbf{i}^{t-1} + \mathbf{R}_{fo} \mathbf{f}^{t-1} + \mathbf{R}_{oo} \mathbf{o}^{t-1}\end{aligned}$$

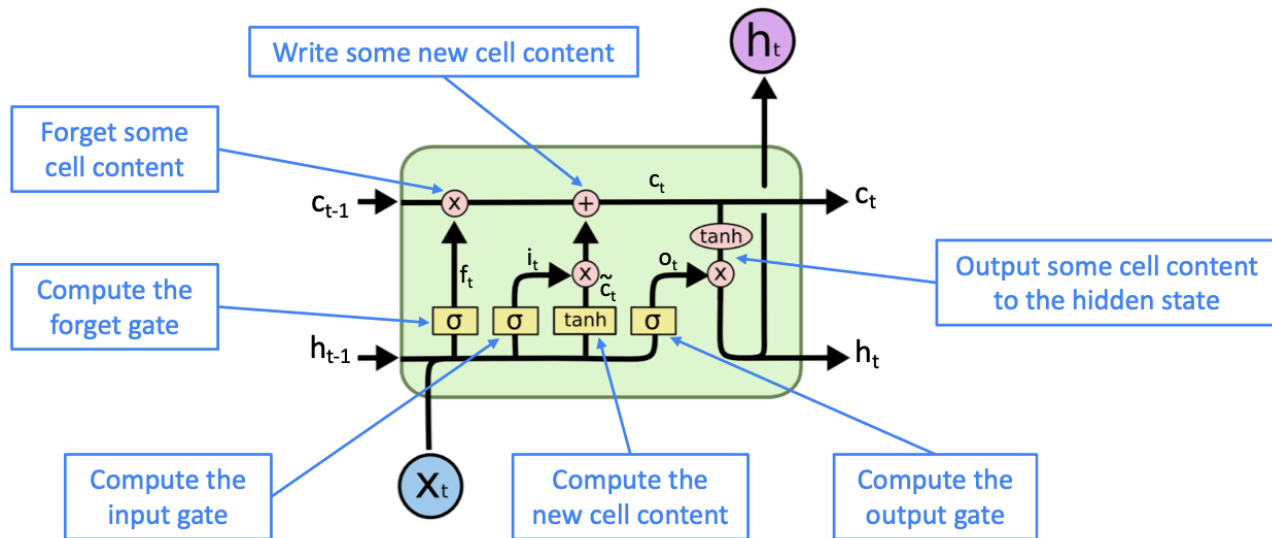
## NP LSTMs

$$\begin{aligned}\bar{\mathbf{i}}^t &= \mathbf{W}_i \mathbf{x}^t + \mathbf{R}_i \mathbf{y}^{t-1} + \mathbf{b}_i \\ \bar{\mathbf{f}}^t &= \mathbf{W}_f \mathbf{x}^t + \mathbf{R}_f \mathbf{y}^{t-1} + \mathbf{b}_f \\ \bar{\mathbf{o}}^t &= \mathbf{W}_o \mathbf{x}^t + \mathbf{R}_o \mathbf{y}^{t-1} + \mathbf{b}_o\end{aligned}$$

# NP LSTMs



# NP LSTMs








# 5. Results and Discussion


It's result time...



# Comparison of LSTM Variants (whole)



# Comparison of LSTM Variants (top 10%)



# Comparison of LSTM Variants

- Removing output activation function and forget gate, drastically reduce LSTM performance.
- CIFG slightly improved performance on Music Modelling, and NP did so in Handwriting Recognition.



# Comparison of LSTM Variants

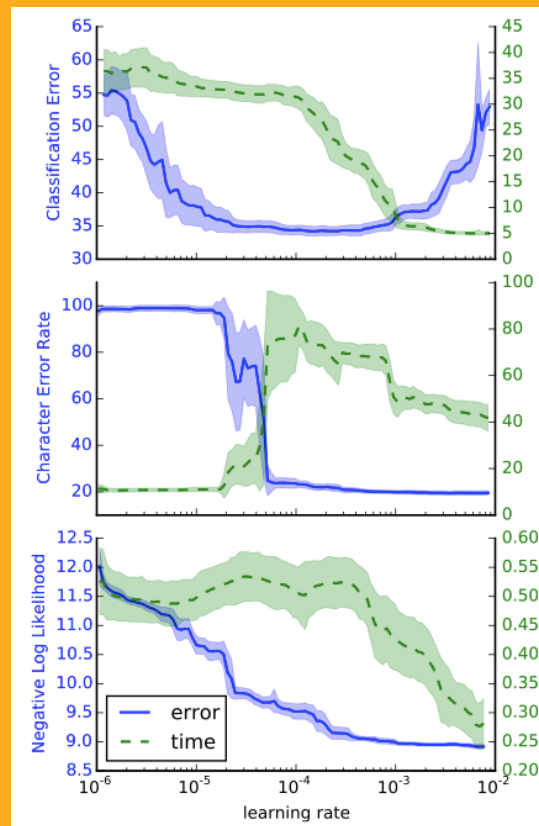
- FGR increases model complexity due to large no. of parameters, and reduces performance on Music Modelling.
- NIG, NOG, NIAF reduce model performance on speech and handwriting recognition significantly, but lead to a slight increase in Music Modelling.

# Impact of Hyper-Parameters

- Hyper-Parameters affect Model performance.
- We discuss how each hyper-parameter affects model performance using the fANOVA framework.
- In fANOVA, we vary one hyper-parameter, while averaging over all others using regression trees.

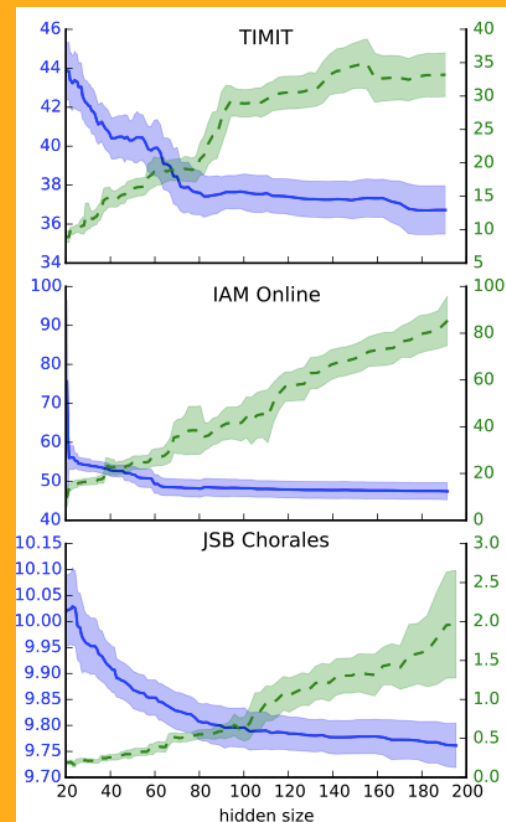
# Learning Rate

- For each dataset, there is large basin of good learning rates inside of which performance doesn't vary much.
- Start with a high value, divide by 10 recursively, until performance stops increasing.



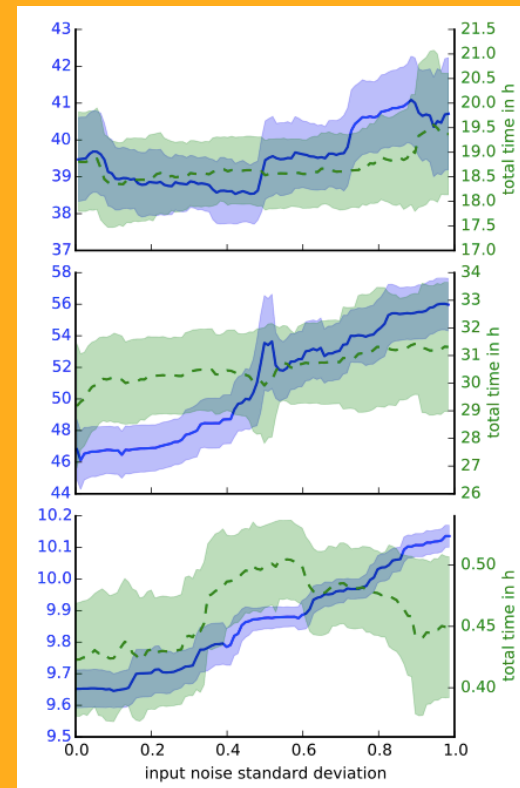
# Hidden Layer Size

- As hidden layer size increases, model performance increases but with diminishing returns.
- After a point of time, the marginal time increases a lot, with very low increases in model performance.



# Input Size


- Additive Gaussian Noise acts as a traditional regularizer for LSTMs just like other Neural Networks.
- But, in the case of LSTMs, it almost always hurts model performance, while increasing training time at the same time.



# Momentum

- Momentum doesn't affect Model performance by more than 1% of the variance of the test set performance.
- This is probably because we scaled the learning rate by a factor of  $(1 - \text{momentum})$ .

# Analysis of Variance





# 6. Conclusion

But what did we learn?





- Vanilla LSTMs perform reasonably well on various datasets.
- None of the 8 variants improved model performance significantly. However, NP and CIFG simplified the model and reduced computational complexity significantly.






- The Forget Gate, and the Output Activation Function are the most critical components of the LSTM.
- Output activation function is used to take care of the unbounded cell state.



“

- Learning Rate is the most important Hyper-Parameter.
  - Hyper-Parameters are almost independent of each other in LSTMs.
- 
- 